

# Taurscribe Architecture Guide for Beginners

**Perfect for:** Complete beginners to programming, Rust newcomers, or anyone curious about how speech recognition works!

**Goal:** Understand how Taurscribe works through simple explanations, fun analogies, and visual diagrams.

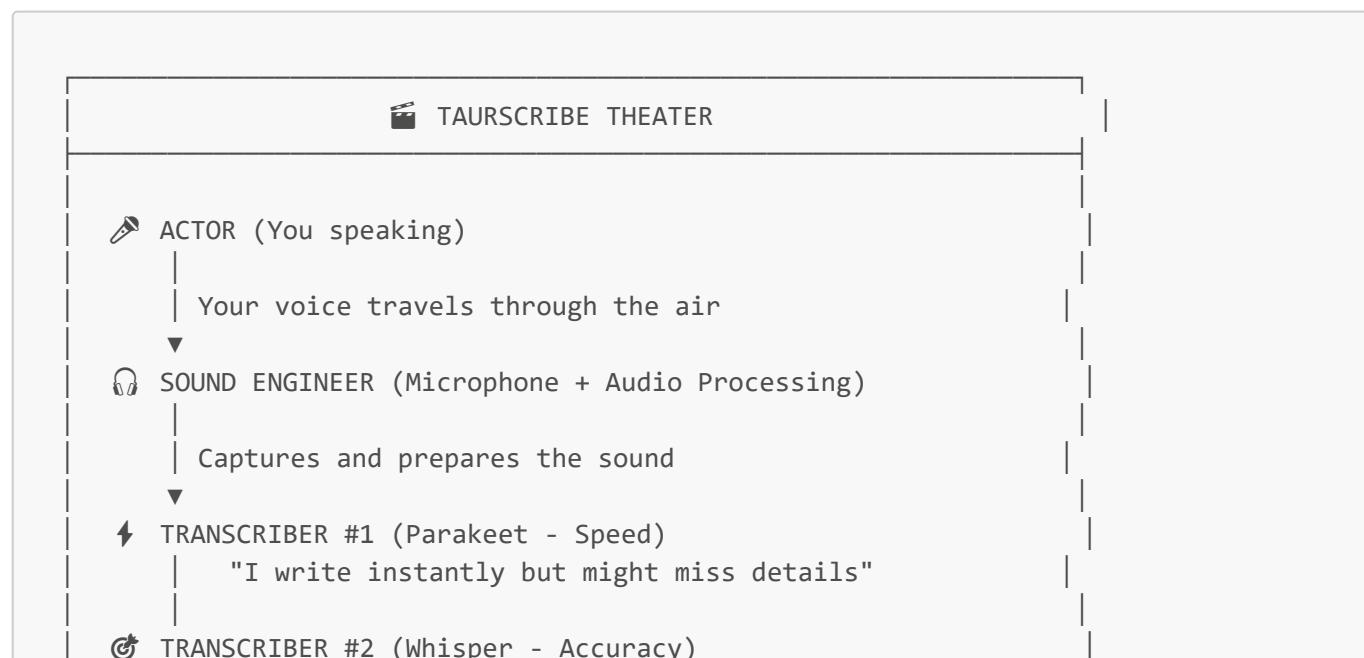
## Table of Contents

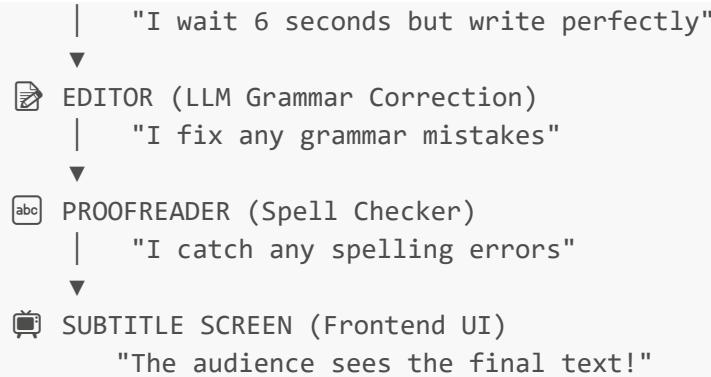
1. [What is Taurscribe?](#)
2. [The Big Picture](#)
3.  [Platform Support & Hardware Acceleration](#)
4.  [Audio Processing: Whisper vs Parakeet](#)
5.  [Voice Activity Detection \(VAD\)](#)
6.  [LLM Integration: Grammar Correction](#)
7.  [Spell Checking](#)
8.  [Model Downloads](#)
9. [Rust Basics You Need to Know](#)
10. [Complete Flow: Start to Finish](#)
11.  [Module Architecture](#)
12.  [Deep Dives: How the Tricky Code Actually Works](#)
13. [File & Function Reference](#)
14. [Common Beginner Questions](#)

## What is Taurscribe?

### Movie Theater Analogy

Imagine Taurscribe is like a **movie theater with live subtitles**:





Taurscribe is a **desktop application** that listens to your voice and magically turns it into text using artificial intelligence!

### Technology Stack (in plain English):

- **Frontend:** React + TypeScript (the pretty buttons and screens you see)
- **Backend:** Rust + Tauri (the super-fast engine that does all the hard work)
- **AI Engines:** Two powerful brains to choose from:
  - **Whisper AI** - Very accurate, great for all situations
  - **Parakeet Nemotron** - Lightning fast, optimized for real-time streaming
- **Post-Processing:**
  - **LLM** - Grammar & style correction with fine-tuned Qwen 2.5 0.5B (GGUF)
  - **Spell Check** - Catch any spelling mistakes with SymSpell

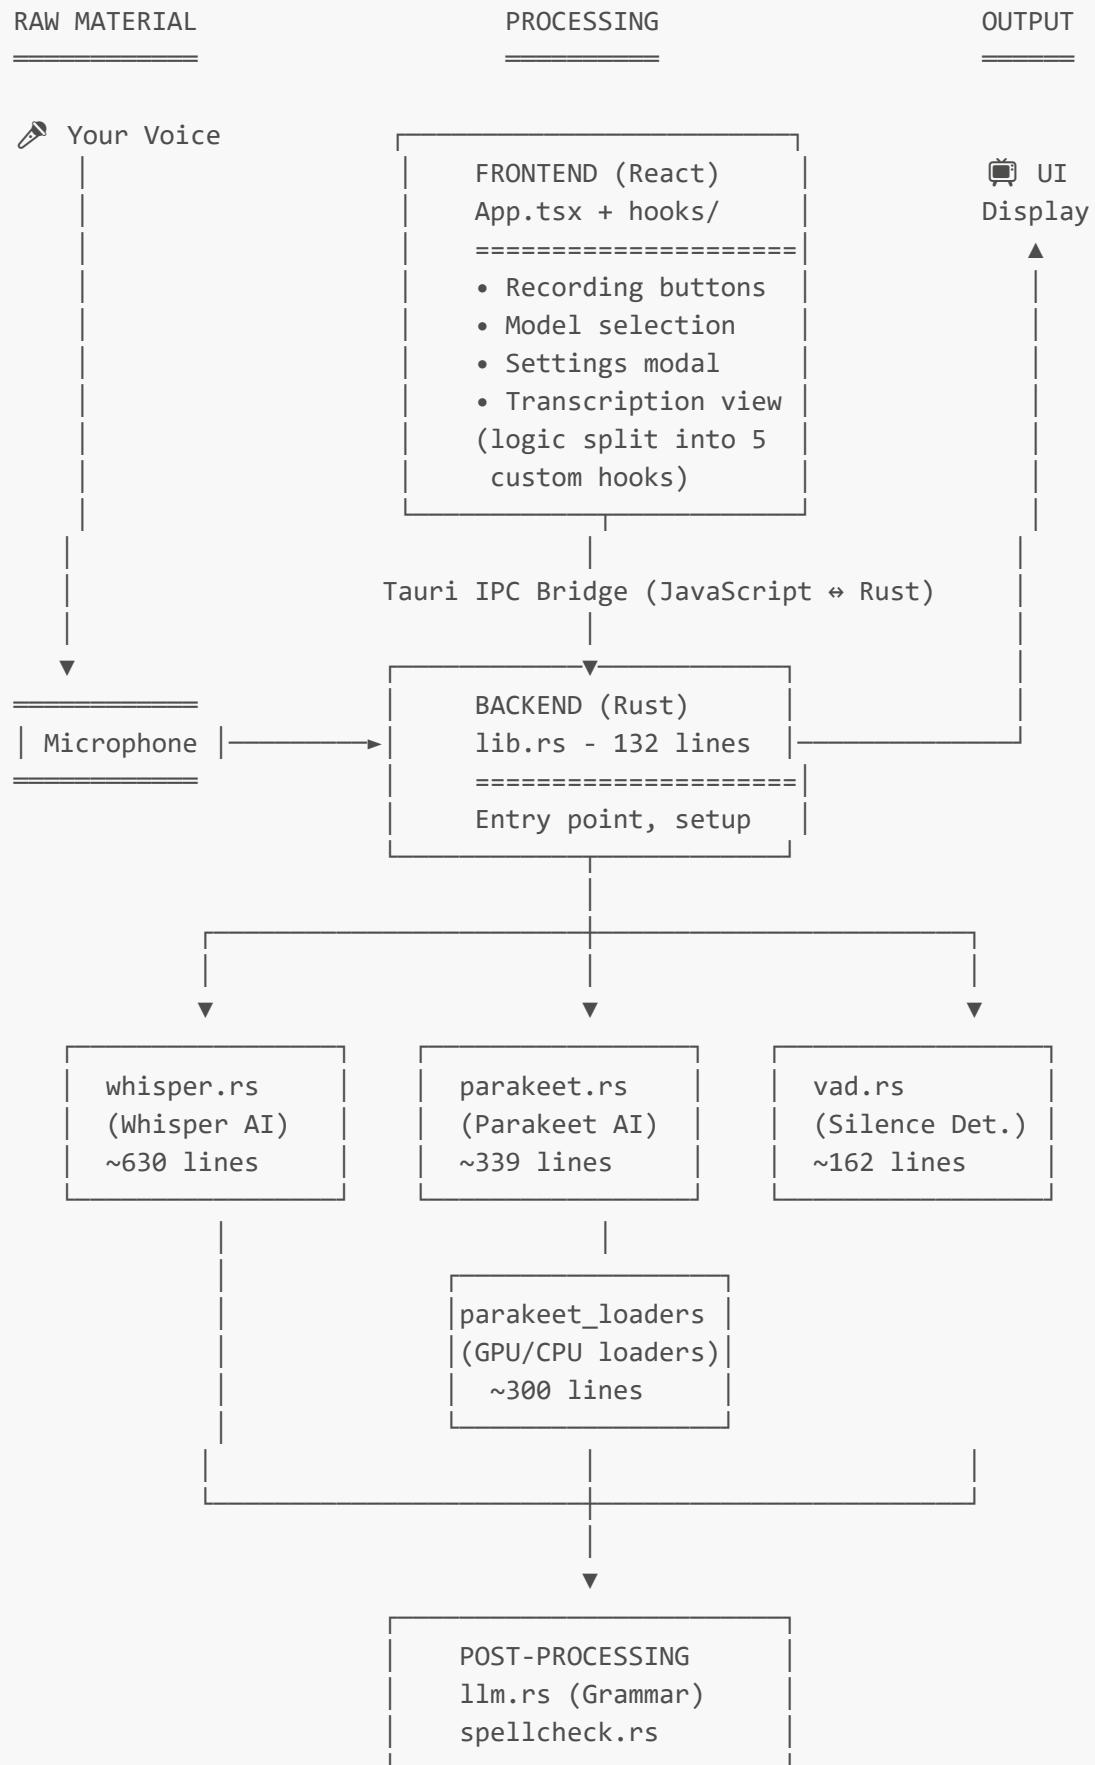
### Key Features:

- Real-time transcription while you speak (see words appear as you talk!)
- High-quality final transcript when you stop
- GPU acceleration for blazing speed (uses your graphics card!)
- Two AI engines to choose from (Whisper or Parakeet)
- Multiple models for each engine (pick small & fast or large & accurate)
- Voice Activity Detection (automatically skips silence)
- Grammar & style correction with local fine-tuned LLM (CPU or GPU)
- Spell checking for final polish
- Model download manager (download models from within the app)
- Global hotkey (Ctrl+Win) works from any application

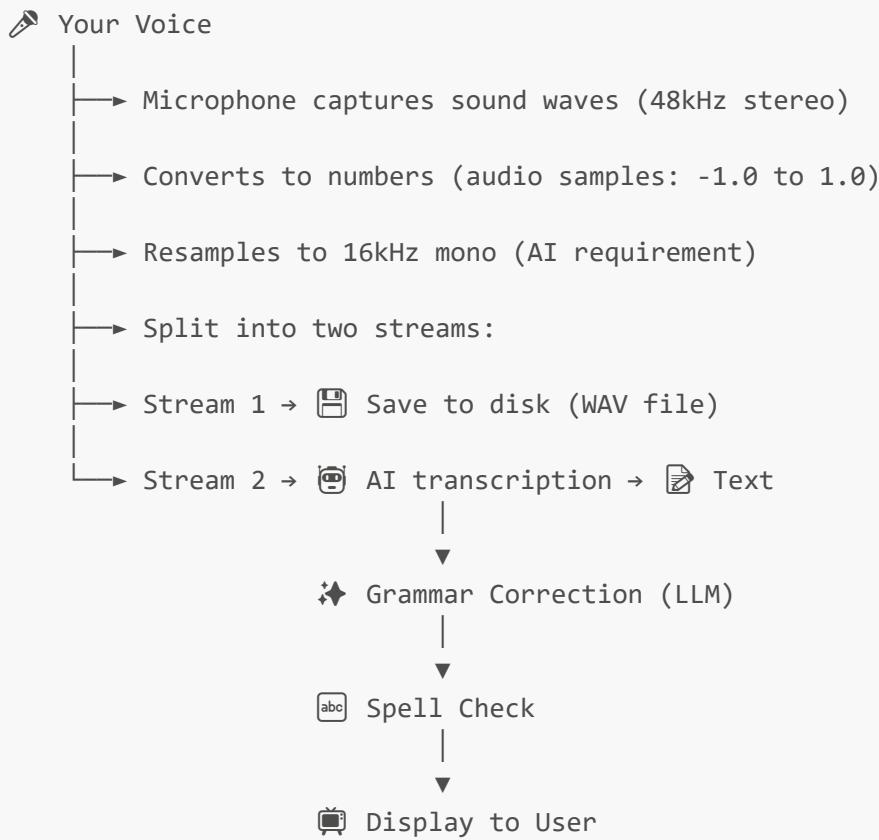
## The Big Picture

### Factory Analogy

Think of Taurscribe as a **speech-to-text factory**:



## Simple Data Flow



### ⚠ Gotcha: Why Two Audio Streams?

**Common Mistake:** Beginners often ask "Why not just use one stream?"

**Answer:** The WAV file is saved in **original quality** (48kHz stereo) while the AI needs **processed audio** (16kHz mono). If we only kept the processed version, we'd lose quality. By saving the original, you can:

- Re-transcribe with different settings later
- Share the original recording
- Use it for other purposes

## 💻 Platform Support & Hardware Acceleration

### 🚗 Car Engine Analogy

Think of hardware acceleration like **different car engines**:

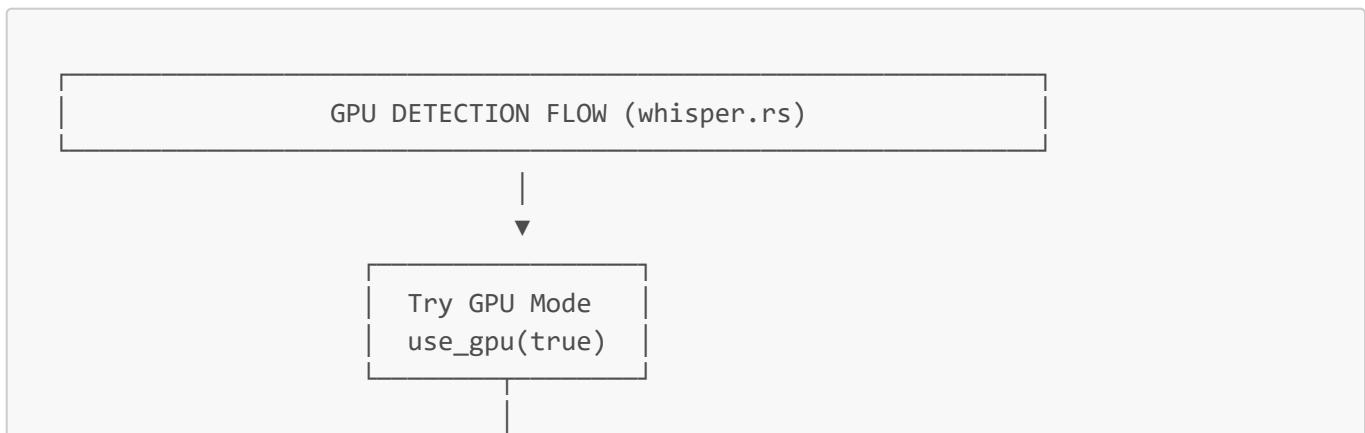
🏎️ ACCELERATION COMPARISON	
⚡ CUDA (NVIDIA GPU)	= Tesla Electric (0-60 in 2s)
Fastest when available, requires NVIDIA	

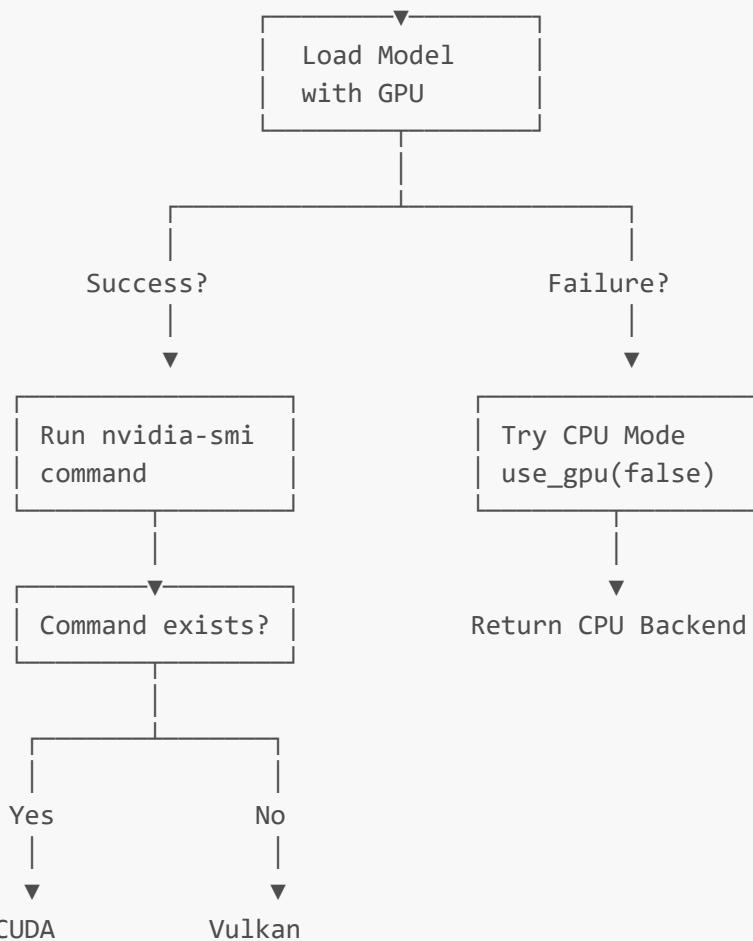
- ⚠️ Vulkan (Any GPU) = Sports Car (0-60 in 4s)  
Good speed, works with AMD/Intel too
- ☒ DirectML (Windows) = Modern Sedan (0-60 in 5s)  
Windows universal, works with NPUs
- ⌚ CoreML (Apple) = BMW Electric (0-60 in 3s)  
Mac-optimized, uses Neural Engine
- ✖️ XNNPACK (CPU) = Economy Car (0-60 in 8s)  
Works everywhere, uses SIMD
- 🏎️ Pure CPU = Bicycle (0-60 in... eventually)  
Always available as fallback

## 📊 Platform Matrix

Platform	Whisper Acceleration	Parakeet Acceleration	Best Use Case
<b>Windows x64 + NVIDIA</b>	CUDA + Vulkan	CUDA + TensorRT	★★★★★ Gaming PCs
<b>Windows x64 + AMD</b>	Vulkan	DirectML	★★★★ AMD systems
<b>Windows ARM64</b>	CPU	DirectML (NPU)	★★★★ Snapdragon laptops
<b>macOS Apple Silicon</b>	Metal	CoreML	★★★★★ MacBook M1/M2/M3
<b>macOS Intel</b>	CPU	XNNPACK	★★★★ Older MacBooks
<b>Linux x64 + NVIDIA</b>	CUDA + Vulkan	CUDA + TensorRT	★★★★★ Linux workstations
<b>Linux ARM64</b>	CPU	XNNPACK	★★★★ Raspberry Pi

## ⌚ How GPU Detection Works





**⚠ Gotcha: CUDA Requires nvidia-smi**

**Common Mistake:** "I have an NVIDIA GPU but it's using Vulkan!"

**Solution:** Make sure NVIDIA drivers are properly installed. The detection runs:

```
std::process::Command::new("nvidia-smi").output()
```

If this fails, Taurscribe assumes Vulkan is available instead.

## 🎙️ Audio Processing: Whisper vs Parakeet

### 🍕 Pizza Delivery Analogy

🎧 AUDIO PROCESSING COMPARISON

- WHISPER AI = Traditional Pizza Delivery
  - Waits for full order (6 seconds of audio)
  - Checks if pizza is worth delivering (VAD check)

- Delivers high-quality pizza (accurate transcription)
- Latency: 6+ seconds

⚡ PARAKEET = Speed Delivery Service

- Delivers slices as they're ready (0.56s chunks)
- No quality check (skips VAD for speed)
- Words appear almost instantly
- Latency: ~0.6 seconds

## 📊 Technical Comparison

Feature	Whisper AI	Parakeet Nemotron
<b>Chunk Size</b>	6.0 seconds (96,000 samples)	0.56 seconds (8,960 samples)
<b>Latency</b>	~6.15 seconds	~0.635 seconds
<b>VAD</b>	<input checked="" type="checkbox"/> Yes (energy-based)	✗ No (speed priority)
<b>Context</b>	Manual (we provide previous text)	Automatic (built-in state via <code>m.reset()</code> )
<b>GPU Support</b>	CUDA, Vulkan, CPU	CUDA, DirectML, CPU
<b>Model Format</b>	GGML (.bin files)	ONNX (.onnx files)
<b>Accuracy</b>	Excellent (95-98%)	Very Good (92-96%)
<b>Best For</b>	Meetings, lectures	Live streaming, gaming

## 🎙 Whisper Processing Pipeline

### WHISPER PIPELINE

#### STEP 1: 🎤 MICROPHONE CAPTURE

Raw Audio: 48,000 samples/second, Stereo, Float32  
 Example: [0.01, -0.02, 0.03, -0.01, 0.04, ...]

#### STEP 2: 🎛 CONVERT TO MONO

Stereo [L1, R1, L2, R2] → Mono [(L1+R1)/2, (L2+R2)/2]  
 Why? AI models expect single-channel audio

#### STEP 3: 🎵 RESAMPLE (48kHz → 16kHz)

Uses `rubato` SincFixedIn resampler (high quality)  
Why? Whisper was trained on 16kHz audio

↓  
STEP 4:  BUFFER INTO 6-SECOND CHUNKS

Accumulate until: `buffer.len() >= 96,000 samples`  
Then: Extract chunk, continue buffering

↓  
STEP 5:  VAD CHECK (Voice Activity Detection)

Calculate RMS (Root Mean Square) "loudness"  
If RMS < 0.005 → Skip (silence)  
If RMS > 0.005 → Process (speech detected)

↓  
STEP 6:  WHISPER AI TRANSCRIPTION

`model.forward(audio_chunk) → "Hello world"`  
Processing time: ~150ms on GPU (40x realtime!)

↓  
STEP 7:  CUMULATIVE CONTEXT

Save transcript for next chunk (`last_transcript` field)  
Helps AI understand: "He said" → who is "he"?

↓  
STEP 8:  SEND TO UI

`emit("transcription-chunk", { text, method: "Whisper" })`

## ⚠ Gotcha: Why 6-Second Chunks?


**Common Mistake:** "Why not 1-second chunks for faster updates?"

**Answer:**

- Too short (1-2s) → Cuts words mid-sentence → AI "hallucinates" (makes up text)
- Too long (30s+) → High latency → Feels slow
- **6 seconds** → Sweet spot: complete sentences + reasonable latency

## 🔊 Voice Activity Detection (VAD)

### 🌐 Traffic Light Analogy



### 💻 How VAD Works (Energy-Based)

```

// Simplified VAD logic from vad.rs
fn is_speech(audio: &[f32]) -> bool {
    // Calculate RMS (Root Mean Square) - a measure of "loudness"
    let sum_squares: f32 = audio.iter().map(|s| s * s).sum();
    let rms = (sum_squares / audio.len() as f32).sqrt();

    // Compare to threshold
    rms > 0.005 // Returns true if louder than threshold
}
    
```

### 📊 VAD Benefits

Feature	Without VAD	With VAD	Benefit
<b>CPU Load</b>	Constant	Low during pauses	Cooler system
<b>Final Speed</b>	~1000ms	~550ms	<b>45% Faster</b>
<b>Accuracy</b>	May hallucinate	Clean silence	No phantom text

## ⚠ Gotcha: VAD Threshold

**Common Mistake:** "VAD keeps marking my speech as silence!"


**Solution:** The threshold (0.005) might be too high for quiet speakers. You can:

1. Increase microphone volume in system settings
2. Speak closer to the microphone
3. (Advanced) Adjust the threshold in `vad.rs`

---

## 🌐 LLM Integration: Grammar Correction

### 📝 Editor Analogy




### 📁 LLM Files Required

```
tauruscribe-runtime/models/qwen_finetuned_gguf/  
└── model_q4_k_m.gguf    ← Fine-tuned Qwen 2.5 0.5B weights (~400 MB)
```

**Note:** The LLM path is resolved in `llm.rs` → `get_grammar_llm_dir()`:

1. **Hardcoded absolute path** — `GRAMMAR_LLM_PATH` const at the top of `llm.rs` (points to the developer's local machine path; update when deploying)
2. Falls back to `GRAMMAR_LLM_DIR` environment variable
3. Final fallback: `%LOCALAPPDATA%\Taurscribe\models\qwen_finetuned_gguf\`

## ⌚ LLM Processing Flow



## ⌚ Transcription Styles

The LLM supports 6 styles selectable from the **LLM & Grammar** settings tab:

Style	What it does
Auto	Default — clean and natural

Style	What it does
<b>Casual</b>	Relaxed tone, contractions kept
<b>Verbatim</b>	Minimal changes, preserves original phrasing
<b>Enthusiastic</b>	Energetic tone, exclamation marks
<b>Software Dev</b>	Preserves technical terms, camelCase, CLI flags
<b>Professional</b>	Formal grammar, business-ready

## ⚠ Gotcha: LLM KV Cache Must Be Cleared

**Common Mistake:** "The LLM crashes after the second transcription!"


**Answer:** Each new request **must** call `ctx.clear_kv_cache_seq(None)` before filling the batch. Without this, llama.cpp throws a sequence inconsistency error and panics.

## ⚠ Gotcha: LLM Backend Selection

The **Auto / GPU** option sets `n_gpu_layers = 99` (offloads all layers to GPU). If GPU loading fails, it automatically retries with `n_gpu_layers = 0` (CPU only). On macOS, GPU is always forced off regardless of selection.

## 📝 Spell Checking

### abc Dictionary Analogy



▼  
Output: "The quick brown fox"

## 📁 Spell Check Implementation

**File:** [src-tauri/src/spellcheck.rs](#)

**Commands:** [src-tauri/src/commands/spellcheck.rs](#)

The spell checker uses **SymSpell** (frequency-based edit distance):

1. Loads a frequency dictionary ([frequency\\_dictionary\\_en\\_82\\_765.txt](#))
2. Splits input text into words
3. For each unknown word, finds closest matches by edit distance
4. Auto-corrects based on word frequency ranking

**Dictionary location:** [%LOCALAPPDATA%\Taurscribe\models\symspell\](#)


### ⚠ Gotcha: Technical Terms

**Common Mistake:** "It keeps marking my technical terms as misspelled!"

**Solution:** Technical terms (like "ONNX", "CUDA", "API") may not be in the dictionary. SymSpell is conservative — it won't auto-correct a word if no close match exists.

## 📥 Model Downloads

### 📦 Package Delivery Analogy



- File deleted if hash mismatch (corrupted)



#### 5 COMPLETION

- emit("download-progress", { status: "done" })
- Frontend refreshes model list
- Model instantly available for use!

## 📁 Available Models (from `commands/model_registry.rs`)

Model ID	Type	Files	Size
<code>whisper-tiny</code>	Whisper GGML	1 <code>.bin</code>	~75 MB
<code>whisper-tiny-q5_1</code>	Whisper GGML (quantized)	1 <code>.bin</code>	~30 MB
<code>whisper-base</code>	Whisper GGML	1 <code>.bin</code>	~142 MB
<code>whisper-base-en</code>	Whisper GGML	1 <code>.bin</code>	~142 MB
<code>whisper-small</code>	Whisper GGML	1 <code>.bin</code>	~466 MB
<code>whisper-small-en</code>	Whisper GGML	1 <code>.bin</code>	~466 MB
<code>whisper-medium</code>	Whisper GGML	1 <code>.bin</code>	~1.5 GB
<code>whisper-large-v3</code>	Whisper GGML	1 <code>.bin</code>	~2.9 GB
<code>whisper-large-v3-turbo</code>	Whisper GGML	1 <code>.bin</code>	~1.6 GB
<code>parakeet-nemotron</code>	Parakeet ONNX	4 files	~700 MB
<code>qwen2.5-0.5b-instruct</code>	GGUF	1 <code>.gguf</code>	~400 MB
<code>qwen2.5-0.5b-instruct-tokenizer</code>	Tokenizer JSON files	4 files	~2 MB
<code>qwen2.5-0.5b-safetensors</code>	SafeTensors (GPU)	multi-file	~1 GB
<code>symspell-en-82k</code>	Dictionary	1 <code>.txt</code>	~6 MB

## ⚠ Gotcha: Download Verification

**Common Mistake:** "The model downloaded but won't load!"

**Answer:** The download might be corrupted. The downloader:

1. Checks SHA-1 hash after download (hash stored in `model_registry.rs`)
2. Deletes the file if hash doesn't match
3. You'll see an error toast if verification fails

Try re-downloading or check your internet connection.

# Rust Basics You Need to Know

## ❖ Ownership Puzzle Analogy

### ❖ RUST OWNERSHIP RULES

Rule 1: Each value has ONE owner

```
let s1 = String::from("hello");
let s2 = s1; // s1 is MOVED to s2
// println!("{}", s1); ← ERROR! s1 no longer valid
```

Rule 2: When owner goes out of scope, value is dropped

```
{
    let s = String::from("hello");
    // s is valid here
} // s is dropped here (memory freed)
```

Rule 3: You can BORROW with references

```
fn print_length(s: &String) { // Borrows, doesn't own
    println!("{}", s.len());
}
let s = String::from("hello");
print_length(&s); // Borrow
println!("{}", s); // Still valid! ✓
```

## ❖ Quick Reference Table

Concept	Syntax	Example
Variable	<code>let x = 5;</code>	<code>let name = "Rust";</code>
Mutable	<code>let mut x = 5;</code>	<code>let mut counter = 0;</code>
Reference	<code>&amp;x</code>	<code>let ref = &amp;value;</code>
Mutable ref	<code>&amp;mut x</code>	<code>let mut_ref = &amp;mut value;</code>
Option	<code>Option&lt;T&gt;</code>	<code>Some(5) or None</code>
Result	<code>Result&lt;T, E&gt;</code>	<code>Ok(5) or Err("error")</code>
Match	<code>match x { ... }</code>	Pattern matching
If let	<code>if let Some(x) = opt { }</code>	Pattern matching shortcut
Unwrap	<code>x.unwrap()</code>	Get value or panic

Concept	Syntax	Example
Question mark	x?	Propagate error

## 🔒 How Shared State Works (`Arc<Mutex<T>>`)

Taurscribe shares engines (Whisper, Parakeet, LLM) across threads safely:

```
// In state.rs - wrapping the WhisperManager for thread-safe sharing
pub whisper: Arc<Mutex<WhisperManager>>,
//           ^^^  ^^^^^^
//           |      └ Mutual Exclusion: only one thread at a time
//           └ Atomic Reference Count: multiple owners across threads

// In commands/recording.rs - using it from a background thread
let whisper = Arc::clone(&state.whisper);
std::thread::spawn(move || {
    let mut w = whisper.lock().unwrap(); // Lock, then use
    w.transcribe_chunk(&audio)?;
});
```

## ⚠ Gotcha: `unwrap()` is Dangerous!

**Common Mistake:** Using `unwrap()` everywhere

**Problem:** `unwrap()` panics if the value is `None` or `Err`, crashing your app!

**Solution:** Use safer alternatives:

```
// ✗ Bad
let value = maybe.unwrap(); // Crashes if None!

// ✅ Good - provide default
let value = maybe.unwrap_or(0);

// ✅ Good - handle both cases
if let Some(v) = maybe {
    println!("Got: {}", v);
}

// ✅ Good - propagate error
let value = maybe.ok_or("No value")?;
```

---

## Complete Flow: Start to Finish

### 💻 Phase 1: User Clicks "Start Recording"

## FRONTEND (useRecording.ts + App.tsx)

```
handleStartRecording() {
  1. Check engine is loaded (Whisper or Parakeet)
  2. invoke("start_recording") → Backend
  3. Set UI state to "Recording"
  4. Update tray icon via invoke("set_tray_state")
}
```

## BACKEND (commands/recording.rs)

```
pub fn start_recording(state: State<AudioState>) {
  1. Clear engine context (last_transcript = "")
  2. Open default microphone (cpal)
  3. Create WAV file writer (hound)
  4. Create channels: file_tx, transcriber_tx
  5. Spawn writer_thread → saves audio to disk
  6. Spawn transcriber_thread → real-time AI inference
  7. Start audio stream (calls callback every ~10ms)
}
```

## ✍ Phase 2: Audio Capture (Every ~10ms)

## AUDIO CALLBACK (runs on CPAL audio thread)

```
move |data: &[f32], _| {
  // 1. Send raw stereo to file writer
  file_tx.send(data.to_vec()).ok();

  // 2. Convert stereo → mono
  let mono = data.chunks(2)
    .map(|c| (c[0] + c[1]) / 2.0)
    .collect();

  // 3. Send mono to transcription thread
  transcriber_tx.send(mono).ok();
}
```

## ⌚ Phase 3: Transcription Thread Loop

```

TRANSCRIPTION THREAD (background thread)

loop {
    // 1. Receive mono audio samples from channel
    let samples = rx.recv()?;
    // Blocks until data arrives

    // 2. Add to ring buffer
    buffer.extend(samples);

    // 3. Check if buffer is large enough
    if buffer.len() >= chunk_size { // 96k for Whisper
        let chunk = buffer.drain(..chunk_size).collect();

        // 4. [Whisper only] Skip silence with VAD
        if engine == Whisper && !vad.is_speech(&chunk) {
            continue; // skip this chunk
        }

        // 5. Resample to 16kHz
        let resampled = resample_to_16k(&chunk);

        // 6. Transcribe with AI engine
        let text = engine.transcribe_chunk(&resampled)?;

        // 7. Emit live result to frontend
        app.emit("transcription-chunk", &text);
    }
} // Loop ends when channel is dropped (recording stopped)

```

## Phase 4: Stop Recording + Post-Processing

```

STOP RECORDING (Frontend → Backend → Frontend)

FRONTEND (useRecording.ts):
1. invoke("stop_recording") → gets raw transcript back

BACKEND (commands/recording.rs):
2. drop(file_tx) → signals writer thread to finish
3. writer_thread.join() → waits for WAV file to finalize
4. [Whisper] Final pass on full WAV file (higher accuracy)
5. [Parakeet] Returns accumulated session transcript
6. clean_transcript() → fixes spacing, punctuation
7. Returns final text to frontend

FRONTEND post-processing pipeline:

```

8. [if spell check ON] → invoke("correct\_spelling")
9. [if grammar LLM ON] → invoke("correct\_text", { style })
10. invoke("type\_text") → Enigo types text into active window
11. Update UI transcript display
12. Update tray icon back to "Ready"

## ⚠ Gotcha: Channel Closing

**Common Mistake:** "The app hangs when I stop recording!"

### How channels work:

1. `drop(file_tx)` closes the **sending end** of the channel
2. The writer thread's `rx.recv()` returns `Err` when the sender is gone
3. Thread exits its loop and finalizes the WAV file
4. **Without `drop()`**, the thread would block forever waiting for data!

## 📐 Module Architecture

### 📁 Current File Structure (Updated February 2026)

```
Taurscribe/
├── Frontend
│   └── src/
│       ├── App.tsx           # UI assembly + event wiring (~440 lines)
│       ├── App.css          # App-level styling
│       ├── main.tsx         # React entry point
│
│       └── hooks/           # All logic lives here – App.tsx just
assembles
|   |   |   ├── useHeaderStatus.ts    # Transient status ticker messages (~25
lines)
|   |   |   ├── useModels.ts        # Whisper + Parakeet model lists (~55 lines)
|   |   |   ├── usePostProcessing.ts # LLM + SymSpell toggle/auto-load (~94
lines)
|   |   |   ├── useEngineSwitch.ts  # Engine switching + model loading (~193
lines)
|   |   |   └── useRecording.ts    # Record start/stop + post-processing (~197
lines)
|
|   └── components/
|       ├── SettingsModal.tsx      # Modal shell + tab router (~357 lines)
|       ├── SettingsModal.css      # Modal styling
|       └── settings/
|           ├── GeneralTab.tsx    # Spell check toggle tab (~90 lines)
|           ├── DownloadsTab.tsx  # Model download list tab (~120 lines)
|           └── ModelRow.tsx       # Single downloadable model row (~130
lines)
```

```

|   |           └── types.ts          # Shared types + MODELS constant (~125
lines)
|   |
|   └── index.html                  # HTML shell
|
└── Backend (Rust)
    └── src-tauri/
        └── src/
            ├── Core
            │   └── lib.rs          # App entry + module declarations (~132
lines)
            |
            ├── main.rs          # Binary entry point (6 lines)
            └── types.rs          # Shared enums: AppState, ASREngine (~30
lines)
            |
            ├── state.rs          # AudioState struct + new() (~68 lines)
            ├── utils.rs          # get_models_dir(), get_recordings_dir(),
            |                           # clean_transcript() (~64 lines)
            └── audio.rs          # RecordingHandle struct (~24 lines)
            |
            └── Audio & ASR Engines
                └── whisper.rs      # WhisperManager: load, transcribe,
resample
                |
                └── parakeet.rs      # ParakeetManager: Nemotron/CTC/EOU/TDT
lines)          # transcription + model status (~339
lines)
                |
                ├── parakeet_loaders.rs # GPU/CPU loader helpers for each
                |                           # Parakeet model type (~300 lines)
                └── vad.rs             # Energy-based VAD: is_speech(),
                                         # get_speech_timestamps() (~162 lines)
                |
                └── Post-Processing
                    └── llm.rs          # LLMEngine: Qwen 2.5 0.5B GGUF via
                                         # llama-cpp-2, format_transcript() (~343
lines)
                    |
                    └── spellcheck.rs     # SymSpell spell checker (~150 lines)
                    |
                    └── Commands (Tauri IPC)
                        └── commands/
                            ├── mod.rs          # Re-exports all pub commands
                            └── recording.rs      # start_recording, stop_recording,
type_text
                            |
                            └── models.rs          # list_models, switch_model,
init_parakeet,
                            |
                            └── llm.rs              # init_llm, unload_llm, correct_text,
etc.                      # check_llm_status
                            |
                            └── spellcheck.rs      # init_spellcheck, unload_spellcheck,
                                         # correct_spelling,
check_spellcheck_status
                            |
                            └── downloader.rs      # download_model, get_download_status,
                                         # delete_model, verify_model_hash

```

```
SHA1s
    |   |   |   model_registry.rs      # get_model_config(): all model URLs +
    |   |   |   settings.rs          # set_tray_state
    |   |   |   misc.rs             # greet (placeholder)

    |   |   └── System Tray
    |   |       └── tray/
    |   |           └── mod.rs      # setup_tray() + icon switching
    |   |           └── (icon assets)

    |   |   └── Global Hotkeys
    |   |       └── hotkeys/
    |   |           └── mod.rs      # Re-exports start_hotkey_listener
    |   |           └── listener.rs # rdev Ctrl+Win listener (~75 lines)

    |   |   └── File Watcher
    |   |       └── watcher.rs     # notify watcher on models dir,
    |   |                           # emits "models-changed" event (~60
lines)


    |   └── build.rs              # macOS deployment target, CUDA linker paths
    |   └── Cargo.toml            # All Rust dependencies + feature flags

    └── Runtime Assets
        └── taurscribe-runtime/
            └── models/
                └── qwen_finetuned_gguf/ # model_q4_k_m.gguf
                └── parakeet-*/*        # ONNX model folders (dev only)

    └── assets/                  # App icons, tray icons (.png / .icns / .ico)

    └── Documentation
        └── ARCHITECTURE.md      # This file!
        └── README.md
```

## E Module Dependency Diagram



```
dependencies)
```

Rule: Lower modules NEVER depend on higher modules!

Frontend hook dependency order:

useHeaderStatus	← (no deps)
useModels	← useHeaderStatus
usePostProcessing	← useHeaderStatus
useEngineSwitch	← useModels, useHeaderStatus
useRecording	← useEngineSwitch, usePostProcessing, useHeaderStatus
App.tsx	← all hooks

## ⚠ Gotcha: Circular Dependencies

**Common Mistake:** "I added `use crate::commands` to `whisper.rs` and it won't compile!"

**Solution:** Lower-level modules (`whisper.rs`, `llm.rs`) must NEVER import from higher-level modules (`commands/`). Instead:


- Put shared types in `types.rs`
- Put utility functions in `utils.rs`
- Let the higher-level module (commands) import from the lower-level ones

## 🔗 Deep Dives: How the Tricky Code Actually Works

These sections break down the most confusing or "magic-looking" parts of the codebase into the simplest possible explanations. Each example is taken directly from the real code.

### 1 Channels — Threads Talking to Each Other

In `commands/recording.rs`, the code uses **channels** to send audio from the microphone thread to the transcription thread. Think of a channel exactly like a **walkie-talkie**:



- `tx.send(data)` → puts data into the pipe (never blocks)
- `rx.recv()` → takes data OUT (BLOCKS until data arrives)
- `drop(tx)` → closes pipe → `rx.recv()` returns Err → thread exits

### Annotated real code from `commands/recording.rs`:

```
// Step 1: Create TWO channels – one for file writing, one for transcription
let (file_tx, file_rx) = crossbeam_channel::unbounded::<Vec<f32>>();
//      ^^^^^^   ^^^^^^
//      Sender  Receiver
//                                         Any size queue (no limit)

let (transcriber_tx, transcriber_rx) = crossbeam_channel::unbounded::<Vec<f32>>();

// Step 2: Spawn a background thread that owns the RECEIVER end
std::thread::spawn(move || {
//           ^
//           "move" = this thread now OWNS transcriber_rx
loop {
    match transcriber_rx.recv() { // ← BLOCKS here, waiting for audio data
        Ok(samples) => { /* transcribe */ }
        Err(_) => break,           // ← tx was dropped = recording stopped
    }
}
});

// Step 3: Audio callback runs on CPAL's thread (every ~10ms)
let callback = move |data: &[f32], _: &_| {
    file_tx.send(data.to_vec()).ok();           // → file writer thread
    transcriber_tx.send(data.to_vec()).ok();     // → transcription thread
    //
    //           ^
    //           .ok() = ignore send error if receiver is gone
};

// Step 4: When recording stops, drop the sender → threads finish naturally
drop(file_tx);          // ← File writer thread sees Err and exits
drop(transcriber_tx);    // ← Transcription thread sees Err and exits
writer_thread.join().unwrap(); // ← Wait for both to finish cleanly
```

#### Gotcha — Why `move` before the closure?

Without `move`, the closure would borrow `transcriber_rx` by reference. But references can't cross thread boundaries in Rust (the original thread might die first). The `move` keyword transfers **ownership** into the new thread, making it safe.

## [2] `Arc<Mutex<T>>` — Sharing a Resource Between Threads

Taurscribe's AI engines (Whisper, Parakeet, LLM) live in `state.rs` and need to be accessed from *multiple* threads. Here's how that works:

 `Arc<Mutex<T>>` = Thread-Safe Safe Deposit Box

## Arc (Atomic Reference Count)

Imagine a "shared photocopy" of a key.  
You can make as many copies as you need.  
The box is destroyed only when ALL copies are gone.

Original Arc    \_\_\_\_ copy 1 (Thread A: recording command)  
                 └ copy 2 (Thread B: transcription thread)  
                 └ copy 3 (Thread C: stop command)

ref count: 3 → box still alive

## Mutex (Mutual Exclusion)

Only ONE thread can look inside the box at a time.  
Others must wait outside until the door opens.

Annotated real code from `state.rs` and `commands/recording.rs`:

```
// In state.rs – declaring the shared state
pub struct AudioState {
    pub whisper: Arc<Mutex<WhisperManager>>,
    //           ^^^  ^^^^^^
    //           |      └ "One thread at a time" lock
    //           └ "Multiple owners" reference-counted pointer

    pub parakeet: Arc<Mutex<ParakeetManager>>,
    pub llm:      Arc<Mutex<LLMEngine>>,
}

// In commands/recording.rs – using WhisperManager from a background thread
pub fn start_recording(state: tauri::State<AudioState>) {
    // Clone the Arc (cheap! just increments the reference count)
    let whisper_clone = Arc::clone(&state.whisper);
    //                                     ^^^^^^^^^^^^^^^^^^
    //                                     borrow to clone – doesn't move the original
```

```

std::thread::spawn(move || {    // Move the clone INTO the thread

    // Lock the mutex - we now have exclusive access
    let mut w = whisper_clone.lock().unwrap();
    //
    // Returns MutexGuard<WhisperManager>
    // Auto-unlocks when `w` goes out of scope (RAII)

    w.transcribe_chunk(&audio)?; // Use WhisperManager safely

}); // `w` drops here = mutex UNLOCKED = other threads can now use it
}

```


### Gotcha — Deadlock!

What if Thread A locks `whisper`, then tries to lock `llm`, while Thread B has `llm` locked and waits for `whisper`? Both threads wait forever — **deadlock!**

**Rule:** Always lock mutexes in the same order everywhere in the code.

## 3 The Tauri IPC Bridge — JavaScript Calling Rust

This is the "magic" that lets the React UI call Rust functions:



## How a command gets registered — `lib.rs`:

```
// lib.rs – this is like a phone directory: "here are all the functions
//           the frontend is allowed to call"
tauri::Builder::default()
    .invoke_handler(tauri::generate_handler![
        //                                     ^^^^^^^^^^^^^^^^^^
        // Macro that wires up the IPC handler
        commands::start_recording,   // JS can call invoke("start_recording")
        commands::stop_recording,    // JS can call invoke("stop_recording")
        commands::correct_text,      // JS can call invoke("correct_text", {text,
style})                                // ...
                                         // ... etc for every command
    ])
)
```

### ⚠ Gotcha — Naming matters!

The string you pass to `invoke("start_recording")` in JavaScript must **exactly** match the Rust function name. A typo gives a silent runtime error, not a compile error.

## 4 Result<T, E> and the ? Operator — Rust Error Handling

Rust has no exceptions. Instead, functions return `Result<Ok, Err>` — a box that contains *either* a success value or an error:

Result<T, E> = A Box With Two Compartments

Result<String, String>

Compartiment A: Ok(String)  
 "The corrected text result"

Compartiment B: Err(String)  
 "Model not loaded: file not found"

The caller **MUST** check which compartment has data before using it.  
Rust forces this – you literally cannot use the value without checking.

### Three ways to handle a Result:

```
// -----
// WAY 1: match - explicit, handle both cases
// -----
match load_model(path) {
    Ok(model)  => { /* use model */ }
    Err(e)      => { eprintln!("Failed: {}", e); }
}

// -----
// WAY 2: ? operator - short-circuit on error (used EVERYWHERE)
// -----
fn start_recording(state: State<AudioState>) -> Result<(), String> {
    let model = load_model(path)?;
    // ^
    // If load_model returns Err(e), this function IMMEDIATELY
    // returns Err(e) - no need to write the match manually.
    // If load_model returns Ok(m), execution continues with m.

    let text = model.transcribe(&audio)?; // Same pattern - bail on error
    Ok(()) // If we got here, everything worked!
}


// -----
// WAY 3: unwrap_or_else - provide a default value
// -----
let dir = get_models_dir().unwrap_or_else(|_| PathBuf::from("/tmp"));
// ^
// If it fails, use /tmp as fallback
```


#### ⚠ Gotcha — ? only works inside functions that return Result

If you try `let x = something()?;` inside `main()` or a closure that returns `()`, the compiler will complain. Wrap the code in a helper function that returns `Result<_, _>` first.

## 5 Thread Lifetimes — What "Spawning" a Thread Actually Means

Here is a visual timeline of how threads start and stop during a recording session:





### Why `join()` matters:

```
// Without join(): WAV file might be half-written when we return!
writing_thread.join().unwrap();
//           ^^^^^^
//           Blocks (waits) until the writing thread finishes
//           finalizing the WAV file header. THEN we return.

// The WAV format requires the file SIZE in the header.
// The writer thread fixes the header LAST, right before it exits.
// Without join(), we'd return a corrupt WAV file.
```

## ⑥ The Audio Resampling Math

Whisper requires 16,000 samples per second. Your microphone records 48,000 samples per second. Here's what resampling actually does:

### 🎵 AUDIO RESAMPLING: 48kHz → 16kHz

Original (48kHz) – 48,000 numbers per second:  
`[0.01, 0.02, 0.03, 0.04, 0.05, 0.06, 0.07, 0.08, 0.09, ...]`

^—————^—————^—————  
position 0      position 2 (→ kept)    position 4 (→ kept)

After resampling (16kHz) – 16,000 numbers per second:  
`[0.02,            0.05,            0.08, ...]`

- The `rubato` library uses a sinc filter (math magic) to:
- Keep every 3rd sample approximately
  - Blend neighboring samples to avoid aliasing (audio distortion)
  - Result: same audio, just 3x fewer numbers

Ratio:  $16000 / 48000 = 1/3 \rightarrow$  output has 1/3 as many samples

### Why not just take every 3rd sample directly?

That's called "downsampling without anti-aliasing" — it causes high-frequency audio artifacts (ugly distortion). The **rubato** sinc resampler applies a low-pass filter first to prevent this. It's like reducing a photo's resolution properly vs. just deleting every 3rd pixel.

## 7 The ChatML Prompt Format — How the LLM "Understands" Instructions

The LLM (Qwen 2.5) uses a special text format called **ChatML** to understand your instructions:

### CHATML FORMAT – The "Protocol" the LLM Speaks

```
<|im_start|>system  
You are Wispr Flow, an AI that transcribes and polishes speech.  
Style: Professional. Fix grammar. Output ONLY the corrected text.  
<|im_end|>
```

▲ "system" message – sets the AI's personality and rules.  
Think of it like a job description given before work starts.

```
<|im_start|>user  
the quick brown fox jump over the lazy dog  
<|im_end|>
```

▲ "user" turn – this is the raw transcription text we want corrected.

```
<|im_start|>assistant
```

▲ We leave this EMPTY – the model fills in the corrected text here.  
It generates: "The quick brown fox jumps over the lazy dog."  
Then stops when it produces <|im\_end|>

Annotated code from **llm.rs**:

```

fn build_chatml_prompt(text: &str, style: &str) -> String {
    format!(
        // The system message – tells the AI what personality to have
        "<|im_start|>system\n\
         You are Wispr Flow, an AI assistant that transcribes speech.\n\
         Style: {style}.\n\
         Output ONLY the corrected text. No explanations.\n\
         <|im_end|>\n\
         \
         <|im_start|>user\n\
         {text}\n\
         <|im_end|>\n\
         \
         <|im_start|>assistant\n",
        // ^ style = style,
        text = text,
    )
}

// During inference, stop generating when we see the end token:
if token == eos_token || decoded.contains("<|im_end|>") {
    break; // LLM is done! Collect what we have.
}

```

## ⑧ Closures — Anonymous Functions ("functions without a name")

Closures appear **everywhere** in Taurscribe. They look confusing at first:

```

// A normal named function:
fn add_one(x: i32) -> i32 {
    x + 1
}

// A closure (same thing, but inline and anonymous):
let add_one = |x: i32| x + 1;
//          ^^^      ^
//  Parameters      Body (no curly braces needed for one expression)

// Multi-line closure:
let process = |data: Vec<f32>| {
    let mono = convert_to_mono(&data);
    resample_to_16k(&mono)
};

```

**The audio callback is a closure capturing variables from the outer scope:**

```
// These variables are declared OUTSIDE the closure:
let file_tx      = /* channel sender */;
let transcriber_tx = /* channel sender */;

// The closure is passed to CPAL as the audio callback.
// It "captures" file_tx and transcriber_tx from the surrounding scope.
let callback = move |data: &[f32], _info: &cpal::InputCallbackInfo| {
    //           ^^^^
    //           Moves captured variables INTO the closure
    //           (transfers ownership – outer scope can no longer use them)

    file_tx.send(data.to_vec()).ok();
    // ^^^^^^^^
    // file_tx was "moved in" above – the closure now owns it

    transcriber_tx.send(data.to_vec()).ok();
};

// CPAL calls this closure every ~10ms on its internal audio thread
let stream = device.build_input_stream(&config, callback, err_fn, None);
```

## ⑨ Iterators and Chaining — Reading the "Fluent" Style

Rust loves chaining iterator methods. Here's how to read them:

```
// Converting stereo [L, R, L, R, ...] to mono [(L+R)/2, ...]
let mono: Vec<f32> = stereo_samples
    .chunks(2)                      // Step 1: Group into pairs [L,R], [L,R], ...
    .map(|chunk| {                  // Step 2: For each pair...
        let left  = chunk[0];       //           get left channel sample
        let right = chunk.get(1).copied().unwrap_or(left); // right (or left if
missing)
        (left + right) / 2.0         //           average them → one mono sample
    })
    .collect();                     // Step 3: Gather all results into Vec<f32>
// ^^^^^^^^^^
// Iterators are LAZY – nothing runs until collect() is called!
```

### 🔗 ITERATOR CHAIN VISUALIZATION

```
Input: [0.1, 0.3, 0.2, 0.4, 0.5, 0.7]
(stereo: left=0.1, right=0.3, left=0.2, right=0.4, ...)

.chunks(2)  →  [0.1, 0.3]  [0.2, 0.4]  [0.5, 0.7]
               |          |          |
.map(avg)   →  0.2        0.3        0.6
```

```
.collect() → [0.2, 0.3, 0.6] ← mono output!
```

## 10 Option<T> — When Something Might Not Exist

Many things in Taurscribe might not exist yet: the loaded model, an active recording, a found word. `Option<T>` represents "maybe a value, maybe nothing":

💡 `Option<T>` = A Box That Might Be Empty

`Some(value)` ← The box HAS something inside  
`None` ← The box is EMPTY

Real example in `state.rs`:

```
pub model: Option<WhisperModel>
//           ^^^^^^
//           model might not be loaded yet!
```

WRONG – crashes if `None`:

```
let m = state.model.unwrap(); // ✗ panics if no model loaded
```

RIGHT – check first:

```
if let Some(m) = &state.model { // ✓ safe
    m.transcribe(&audio)?;
} else {
    return Err("No model loaded".into());
}
```

## 11 The VAD Math Explained Simply (RMS)

The VAD (Voice Activity Detection) uses a formula called **RMS (Root Mean Square)**. Here's what it means in plain English:

💡 RMS FORMULA – STEP BY STEP

Input audio chunk: `[0.01, -0.02, 0.03, -0.01]`

Step 1: SQUARE every sample (makes all numbers positive)  
 $[0.01^2, 0.02^2, 0.03^2, 0.01^2]$

```
[0.0001, 0.0004, 0.0009, 0.0001]

Step 2: AVERAGE the squares (sum / count)
(0.0001 + 0.0004 + 0.0009 + 0.0001) / 4 = 0.000375

Step 3: SQUARE ROOT (undo the squaring from step 1)
√0.000375 ≈ 0.019

Result: RMS = 0.019
Threshold: 0.005
0.019 > 0.005 →  SPEECH DETECTED

Intuition: RMS = "average loudness" of the audio chunk
• Loud speech → high RMS (e.g., 0.05–0.2)
• Quiet room → low RMS (e.g., 0.001–0.003)
• Threshold 0.005 = the dividing line between speech and silence
```

## File & Function Reference

### 🔍 Quick Lookup Table

I want to...	Go to	Function/Item
Add a new Tauri command	<code>commands/*.rs</code>	Add <code>#[tauri::command]</code> fn + register in <code>lib.rs</code>
Change recording behavior	<code>commands/recording.rs</code>	<code>start_recording()</code> , <code>stop_recording()</code>
Modify Whisper logic	<code>whisper.rs</code>	<code>transcribe_chunk()</code> , <code>WhisperManager::new()</code>
Modify Parakeet transcription	<code>parakeet.rs</code>	<code>transcribe_chunk()</code> , <code>initialize()</code>
Change how Parakeet loads GPU/CPU	<code>parakeet_loaders.rs</code>	<code>init_*_gpu()</code> , <code>init_*_cpu()</code>
Add a new downloadable model	<code>commands/model_registry.rs</code>	Add entry to <code>get_model_config()</code>
Add model to the downloads UI	<code>src/components/settings/types.ts</code>	Add entry to <code>MODELS</code> array
Change download/verify logic	<code>commands/downloader.rs</code>	<code>download_model()</code> , <code>verify_model_hash()</code>

I want to...	Go to	Function/Item
Change LLM prompt or style	llm.rs	format_transcript()
Change LLM inference params	llm.rs	run_with_options()
Change spell check	spellcheck.rs	Correction logic
Modify tray icon/behavior	tray/mod.rs	setup_tray()
Change global hotkey	hotkeys/listener.rs	Modify key match arms
Add shared enum/struct	types.rs	Define struct/enum
Add utility function	utils.rs	Create pub fn
Change UI recording logic	src/hooks/useRecording.ts	handleStartRecording(), handleStopRecording()
Change engine switching UI	src/hooks/useEngineSwitch.ts	handleSwitchToWhisper(), handleSwitchToParakeet()
Change LLM/spell UI toggles	src/hooks/usePostProcessing.ts	Toggle + load/unload logic
Change settings tabs	src/components/SettingsModal.tsx	renderContent(), tab list
Modify General settings tab	src/components/settings/GeneralTab.tsx	Spell check toggle UI
Modify Downloads tab	src/components/settings/DownloadsTab.tsx	Model list + ModelRow

## >All Tauri Commands (as of February 2026)

```
// From lib.rs invoke_handler - matches tauri::generate_handler! exactly:

// 🔑 Misc
commands::greet, // Test/greeting placeholder

// 🗣 Recording
commands::start_recording, // Start mic + real-time transcription
commands::stop_recording, // Stop + final transcript + post-process
commands::type_text, // Type text via Enigo keyboard injection

// 💬 Whisper model management
commands::list_models, // List downloaded Whisper .bin files
commands::get_current_model, // Get active Whisper model name
```

```

commands::switch_model,                                // Load a different Whisper model

// ⚡ Parakeet model management
commands::list_parakeet_models,                      // List Parakeet models + their status
commands::init_parakeet,                            // Initialize a Parakeet model (GPU/CPU)
commands::get_parakeet_status,                      // Check if Parakeet is loaded + which model

// 🎛 Engine switching
commands::set_active_engine,                        // Switch between Whisper / Parakeet
commands::get_active_engine,                         // Get the currently active engine
commands::get_backend_info,                          // Get GPU backend info string

// 📽 System tray
commands::set_tray_state,                           // Update tray icon
(Ready/Recording/Processing)

// ✨ LLM grammar correction
commands::init_llm,                               // Load Owen GGUF model (GPU or CPU)
commands::unload_llm,                             // Unload LLM to free VRAM
commands::run_llm_inference,                      // Raw LLM text generation
commands::check_llm_status,                        // Returns bool: true = loaded, false = not
loaded
commands::correct_text,                           // Format transcript with style via LLM

// 📄 Spell checking
commands::init_spellcheck,                        // Load SymSpell dictionary
commands::unload_spellcheck,                      // Unload spell checker
commands::check_spellcheck_status,                // Check if spell checker is loaded
commands::correct_spelling,                        // Run SymSpell correction on text

// 📾 Download manager
commands::download_model,                         // Stream download from Hugging Face
commands::get_download_status,                    // Check downloaded/verified status per model
commands::delete_model,                           // Delete model file(s) from disk
commands::verify_model_hash,                     // Verify SHA-1 integrity of model file

```

## Common Beginner Questions

Q1: Why are there two transcription engines?

**Answer:** Different use cases need different trade-offs:

- **Whisper** — Higher accuracy, 6-second latency → Best for dictation, meetings
- **Parakeet** — Lower latency (~0.6s), slightly less accurate → Best for real-time streaming

Q2: Can I use this for other languages?

Whisper supports 99 languages — just speak and it auto-detects. Parakeet is English-only (NVIDIA Nemotron model).

Q3: How much RAM does this use?

Component	RAM Usage
Whisper tiny	~100 MB
Whisper base	~200 MB
Whisper large-v3	~3 GB
Parakeet Nemotron	~500 MB
Qwen LLM (Q4_K_M)	~400 MB
Audio buffer	~10 MB

LLM and Spell Checker are **not loaded at startup** — only when you enable them.

Q4: Why does the first transcription take longer?

**Answer:** GPU "warm-up"! The first run compiles CUDA/Vulkan shader kernels. Taurscribe optionally runs a warm-up pass during model initialization to hide this delay from the user.

Q5: What if my recording crashes mid-session?

#### Safety features:

1. WAV file is written continuously stream → disk (you don't lose audio)
2. File saved to: `%LOCALAPPDATA%\Taurscribe\temp\`
3. You can manually re-transcribe the WAV with any tool

Q6: Where do downloaded models go?

All models land in `%LOCALAPPDATA%\Taurscribe\models\`:

```
models/
├── ggml-tiny.bin           ← Whisper models
├── ggml-base.en.bin
└── parakeet-nemotron/
    ├── encoder.onnx
    └── decoder.onnx
└── qwen_finetuned_gguf/
    └── model_q4_k_m.gguf   ← Grammar LLM
└── symspell/                ← Spell check dictionary
    └── frequency_dictionary_en_82_765.txt
```

Q7: How does the global hotkey work?

`hotkeys/listener.rs` spawns a background thread that uses `rdev::listen()` to capture **every** key event system-wide. When both `Ctrl + Win (Meta)` are held:

- Sends `hotkey-start-recording` event → Frontend starts recording
- On key release → Sends `hotkey-stop-recording` → Frontend stops recording

# Conclusion

Taurscribe demonstrates modern Rust practices in a real production app:

- Ownership** — Threads take ownership of data they need
- Borrowing** — Functions borrow without taking ownership
- Concurrency** — Multiple threads work safely in parallel
- Error Handling** — `Result`, `? operator`, `anyhow` for safety
- Modularity** — Clean separation into focused modules after refactoring

## Architecture Benefits:

Feature	Benefit
Separate threads	UI never freezes during AI inference
Crossbeam channels	Safe, backpressure-aware thread communication
<code>Arc&lt;Mutex&lt;T&gt;&gt;</code>	Shared engine state protection
Two AI engines	User picks speed OR accuracy
GPU acceleration	12–60× faster than CPU-only
<code>commands/</code> split	Each command file has one clear responsibility
<code>model_registry.rs</code>	Single source of truth for all model configs
On-demand loading	Parakeet + LLM don't use memory until needed

**Key Takeaway:** Rust's strict compiler prevents entire categories of bugs (data races, null pointer crashes, use-after-free). Once your code compiles, it usually works correctly!

---

## Next Steps

### To learn more Rust:

1. [The Rust Book](#) — Official, comprehensive
2. [Rust By Example](#) — Learn by doing
3. [Rustlings](#) — Interactive exercises

### To extend Taurscribe:

1. Add a new Whisper or Parakeet model variant (edit `model_registry.rs` + `types.ts`)
2. Add a new transcription style to the LLM (edit `format_transcript()` + the style dropdown)
3. Implement speaker diarization (who's speaking)
4. Add export formats (SRT, VTT, plain TXT)
5. Replace energy-based VAD with Silero neural VAD for higher accuracy

**Questions?** Review this guide, check code comments, or explore the Rust documentation!